

STAT2401: Analysis of Experiments

Simple Linear Regression

John W. Lau

Dept. of Math. & Stat.

April 13, 2019

Original-Authors: Ed Cripps, Kevin Murray, and Russel John
Special Thanks: Kenyon Ng

Relationship Between Two Variables

- Measure a pair of variables on a sample of experimental units or subjects:
 - height and weight of a random sample of people
 - income of wife and husband in a sample of couples
 - GDP and Debt of a sample of countries
 - response and dose of drug in a sample of patients
- The Variables:

Y — response variable or dependent variable

X — explanatory variable or independent variable

- Two related but distinct questions of interest:
 - Are Y and X related?
 - Can Y be predicted from X ?

Which Variable is Which and Does It Matter?

- Context will usually determine which variable is X and which Y
- Y is the variable of interest and X the variable which is being used in an attempt to explain what might be contributing to the variability in Y
- The variability is assumed to be in the Y 's and the X 's are taken as fixed
- The explanation of Y is conditional on the values of X observed
- Reversing the roles of X and Y does change the nature of the problem and will lead to different results and conclusions

Simple Linear Regression Model

- The model assumes the expected value of Y for a given X is a straight line i.e.

$$E(Y|X) = \mu(Y|X) = \beta_0 + \beta_1 X$$

- Parameters β_0 and β_1 are constants which do not depend on X
- It also assumes the distribution of Y for any X is Normal and that

$$\text{Var}(Y|X) = \sigma^2$$

does not depend on X

- If $\beta_1 = 0$ then Y and X are unrelated, for then the distribution of Y given X does not depend on X
- The problem of establishing a relationship devolves to determining whether $\beta_1 = 0$ is plausible, based on the magnitude of some sample estimate of β_1

Simple Linear Regression

Simple Linear Regression Model of a Sample of (X, Y) Pairs

- For a given set of observations (X_i, Y_i) , $i = 1 \dots n$


$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

where β_0 and β_1 are fixed parameters to be estimated and ϵ_i are independent, identically distributed Normal random variables with mean 0 and constant variance σ^2

- All variability is in the response variable
- The model describes the conditional distribution of Y at the X points observed in the sample
- At each point, we sample a value from a Normal distribution whose mean lies on the straight line determined by β_0 and β_1 and whose variance remains constant.

Simple Linear Regression


Given a scatter plot of the data



Find the straight line which "best" fits the data

Simple Linear Regression


Equation of the Line: $y = \beta_0 + \beta_1 x$



- β_0 is the *intercept*; $y = \beta_0$ when $x = 0$
- β_1 is the *slope*; if x increases 1 unit, y changes by β_1 units

Simple Linear Regression

Least Squares Regression Line of Y on X



- Vertical displacement of points from line are *residuals*
- The sum of the squared residuals is a measure of how close the line goes to the points
- Choose β_0 and β_1 to minimize the sum of the squared residuals

Normal Equations (Optional)

- The sum of the squared residuals (the objective function) to be minimized is

$$\sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i)^2$$

- Differentiating the objective function with respect to β_0 and β_1 and setting both derivatives to 0 gives the Normal equations:

$$\begin{aligned}\sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) &= 0 \\ \sum_{i=1}^n X_i (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) &= 0\end{aligned}$$

This implies

$$\begin{aligned}\hat{\beta}_0 + \hat{\beta}_1 \bar{X} &= \bar{Y} \\ \hat{\beta}_0 + \hat{\beta}_1 \frac{1}{n} \sum_{i=1}^n X_i^2 &= \frac{1}{n} \sum_{i=1}^n X_i Y_i\end{aligned}$$

where $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ and $\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$

Normal Equations

- The first Normal equation gives

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

Substituting this into the second gives

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n X_i(Y_i - \bar{Y})}{\sum_{i=1}^n X_i(X_i - \bar{X})}$$

- Noting that $\sum_{i=1}^n (Y_i - \bar{Y}) = 0 = \sum_{i=1}^n (X_i - \bar{X})$ leads to two more expressions

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{\sum_{i=1}^n (X_i - \bar{X})Y_i}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

Fitted Values and Unbiasedness

- The fitted value for an observation Y_i is the estimate of $E(Y_i)$, namely the point on the least squares regression line at $X = X_i$:

$$\hat{\mu}(Y|X_i) = \hat{\beta}_0 + \hat{\beta}_1 X_i = \bar{Y} - \hat{\beta}_1(X_i - \bar{X})$$

- The latter formulation shows the least squares regression line passes through the point (\bar{X}, \bar{Y})
- From the model, $E(Y_i - \bar{Y}) = \beta_1(X_i - \bar{X})$ and so

$$E(\hat{\beta}_1) = \frac{\sum_{i=1}^n X_i E(Y_i - \bar{Y})}{\sum_{i=1}^n X_i (X_i - \bar{X})} = \beta_1$$

and

$$E(\hat{\beta}_0) = E(\bar{Y}) - E(\hat{\beta}_1)\bar{X} = \beta_0 + \beta_1\bar{X} - \beta_1\bar{X} = \beta_0$$

- The fitted value is thus an unbiased estimate of $E(Y_i)$, since

$$E(\hat{\mu}(Y|X_i)) = \beta_0 + \beta_1 X_i = \mu(Y|X_i)$$

Absolute Magnitude of β_1

- From any of the given formulas for β_1 it is clear that it is not a dimensionless constant
- If Y is measured in kg and X in m, the units of β_1 would be kg/m
- Changing the units of Y to g would increase the absolute magnitude of β_1 a thousandfold
- So the absolute magnitude of β_1 is no measure of the strength of the relationship between Y and X , nor can it be used on its own to determine if Y is related to X
- One way of proceeding is to standardize the X and Y sample values before fitting the least square regression line
- To standardize a sample, subtract the mean of the sample from all observations and then divide them all by the standard deviation of the sample

Simple Linear Regression

Example: Ponds Institute

- Price (per ounce) of 36 facial cleansers and a preference rating for each based on scores from a panel of 90 women oblivious to brand

```
> load("facecleanser.RData")
> str(facecleanser)

'data.frame': 36 obs. of  4 variables:
 $ price   : num  0.85 2.13 1.19 2.75 4 4 2.5 2.46 2.25 2.23 ...
 $ rating  : num  61 60 60 57 57 56 56 56 54 53 ...
 $ Z.price : num  -0.803 0.426 -0.477 1.021 2.221 ...
 $ Z.rating: num   1.58 1.46 1.46 1.1 1.1 ...
 - attr(*, "variable.labels")= Named chr   "" "" "Zscore(PRICE)" "Zscore(RATING)"
 ..- attr(*, "names")= chr   "price" "rating" "zprice" "zrating"
 - attr(*, "codepage")= int 1252
```


The Standardized Price Z.price and Standardized rating Z.rating are given by

```
> facecleanser$Z.price = with(facecleanser, (price - mean(price)) / sd(price))
> facecleanser$Z.rating = with(facecleanser, (rating - mean(rating)) / sd(rating))
```

Simple Linear Regression

Example: Ponds Institute

```
> with(facecleanser, plot(rating~price, pch=16))  
> with(facecleanser, lines(fitted(lm(rating~price))~price))  
> with(facecleanser, plot(Z.rating~Z.price, pch=16))  
> with(facecleanser, lines(fitted(lm(Z.rating~Z.price))~Z.price))
```



Sample Correlation Coefficient

- Let

$$Z_{X_i} = \frac{X_i - \bar{X}}{S_X} \text{ and } Z_{Y_i} = \frac{Y_i - \bar{Y}}{S_Y}$$

be the standardized samples where

$$S_X = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2}, S_Y = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2},$$

$$\bar{X} = \sum_{i=1}^n X_i, \text{ and } \bar{Y} = \sum_{i=1}^n Y_i.$$

- Note $\sum_{i=1}^n (Z_{X_i} - \bar{Z}_X)^2 = \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{S_X^2} = n - 1$ and

$$\bar{Z}_X = \frac{1}{n} \sum_{i=1}^n Z_{X_i} = 0 \text{ and similarly for } Z_{Y_i}$$

- The slope of the least squares regression line of Z_Y on Z_X is called the sample correlation coefficient, written r_{XY} (or R)

$$r_{XY} = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2 \sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

Sample Correlation Coefficient

- Replacing X_i and Y_i by Z_{X_i} and Z_{Y_i} in the formula, that is to say if we consider $Z_{Y_i} = \beta_0 + \beta_1 Z_{X_i} + \epsilon_i$, $\hat{\beta}_1$ gives

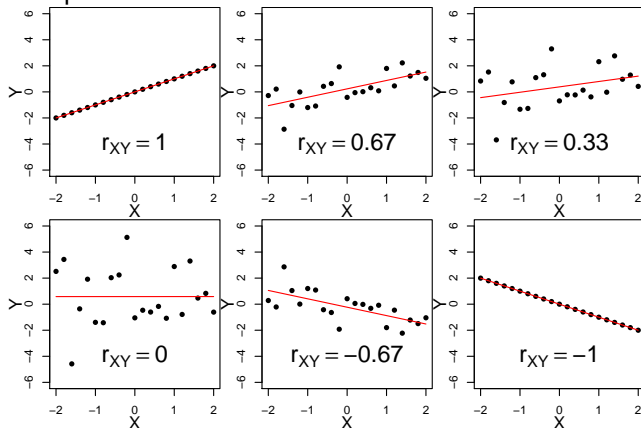
$$\begin{aligned}\hat{\beta}_1 &= \frac{\sum_{i=1}^n (Z_{X_i} - \bar{Z}_X)(Z_{Y_i} - \bar{Z}_Y)}{\sum_{i=1}^n (Z_{X_i} - \bar{Z}_X)^2} \\&= \frac{\sum_{i=1}^n \frac{(X_i - \bar{X})}{s_X} \frac{(Y_i - \bar{Y})}{s_Y}}{n-1} \\&= \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2 \sum_{i=1}^n (Y_i - \bar{Y})^2}} = r_{XY}\end{aligned}$$

- $|r_{XY}| \leq 1$ (Cauchy-Schwarz inequality)
- r_{XY} is +1 or -1 if Y_i and X_i lie in a straight line
- r_{XY} is 0 if Y_i and X_i are not linearly related

Simple Linear Regression

Scatterplots of Standardized X and Y with Various Values of r_{XY}

- Some Scatterplots



r As Sample Estimate of ρ

- The correlation coefficient of two random variables Y and X is

$$\rho_{XY} = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X)\text{Var}(Y)}}$$

where $\text{Cov}(X, Y) = E(X - E(X))(Y - E(Y))$

- ρ is a measure of association between the two variables
- If X and Y are independent then $\rho_{XY} = 0$, but more conditions (like joint Normality) must be imposed on the Y and X for $\rho_{XY} = 0$ to imply independence


In this, both Y and X are variables in their own right with a joint probability distribution

- It is the particular properties of joint Normality which allows the analysis to proceed via the conditional distribution of Y given X

Simple Linear Regression

Components of Total Variability

- What is contributing to the variability in the Y 's?




Large variability about the line and almost no contribution from the X 's

Simple Linear Regression

Components of Total Variability

- Exactly the same variability in the Y 's



Little variability about the line; Y inherits its variability from the X 's

Simple Linear Regression

The Arithmetic of the Decomposition (Optional)

- The fitted value for Y_i given X_i

$$\hat{\beta}_0 + \hat{\beta}_1 X_i$$

- The residual is given by

$$Y_i - (\hat{\beta}_0 + \hat{\beta}_1 X_i)$$

- The difference between Y_i and \bar{Y}

$$\begin{aligned} Y_i - \bar{Y} &= (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) - (\bar{Y} - \hat{\beta}_0 - \hat{\beta}_1 \bar{X}) && \Leftarrow \text{(Total)} \\ &= (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) && \Leftarrow \text{(Residual)} \\ &\quad + (\hat{\beta}_0 + \hat{\beta}_1 X_i - \bar{Y}) && \Leftarrow \text{(Regression)} \\ &= (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) && \Leftarrow \text{(Residual)} \\ &\quad + (\hat{\beta}_0 + \hat{\beta}_1 X_i - (\hat{\beta}_0 + \hat{\beta}_1 \bar{X})) && \Leftarrow \text{(Regression)} \\ &= (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) && \Leftarrow \text{(Residual)} \\ &\quad + \hat{\beta}_1 (X_i - \bar{X}) && \Leftarrow \text{(Regression)} \end{aligned}$$

Simple Linear Regression

The Arithmetic of the Decomposition (Optional)

- Putting "sum of squares" in both sides

$$\underbrace{\sum_{i=1}^n (Y_i - \bar{Y})^2}_{\text{Total Sum of Squares}} = \underbrace{\sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2}_{\text{Residual Sum of Squares}} + \underbrace{\hat{\beta}_1^2 \sum_{i=1}^n (X_i - \bar{X})^2}_{\text{Regression Sum of Squares}}$$


because the cross term is zero, i.e.

$$\begin{aligned} & 2\hat{\beta}_1 \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)(X_i - \bar{X}) \\ &= 2\hat{\beta}_1 \sum_{i=1}^n Y_i(X_i - \bar{X}) - 2\hat{\beta}_1^2 \sum_{i=1}^n X_i(X_i - \bar{X}) \\ &= 2\hat{\beta}_1 \left(\hat{\beta}_1 \times \sum_{i=1}^n X_i(X_i - \bar{X}) \right) - 2\hat{\beta}_1^2 \sum_{i=1}^n X_i(X_i - \bar{X}) = 0 \end{aligned}$$

Simple Linear Regression

Competing Models

- Full Model and Reduced Model



Fitted value: $\hat{\beta}_0 + \hat{\beta}_1 X_i$: 2 parameters
Residual: $Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i$
Residual SS: Residual SS on $n - 2$ df

\bar{Y} : 1 parameter
 $Y_i - \bar{Y}$
Total SS on $n - 1$ df

Extra-Sum-of-Squares Principle:

- Difference in Residual SS (Reduced – Full) = Regression SS
- Difference in Residual DF (Reduced – Full) = 1
- Regression MS/Residual MS (Full) = F -ratio

Simple Linear Regression

Example: Ponds Institute

```
> F1 = lm(rating~price,data=facecleanser)  
> anova(F1)
```

Analysis of Variance Table

Response: rating

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
price	1	214.72	214.723	3.2239	0.08146
Residuals	34	2264.50	66.603		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Example: Ponds Institute

```
• > summary(F1)
```

Call:

```
lm(formula = rating ~ price, data = facecleanser)
```

Residuals:

Min	1Q	Median	3Q	Max
-19.061	-3.607	2.235	4.666	15.268

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	43.711	2.616	16.712	<2e-16 ***
price	2.378	1.325	1.796	0.0815 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.161 on 34 degrees of freedom

Multiple R-squared: 0.08661, Adjusted R-squared: 0.05974

F-statistic: 3.224 on 1 and 34 DF, p-value: 0.08146

Simple Linear Regression

Estimating σ^2 and the SE's

- The Residual MS is an unbiased estimate of σ^2 . Later we will see

$$E\left[\frac{1}{n-2} \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2\right] = \sigma^2$$

- Using $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$ in the last formula for $\hat{\beta}_1$ gives

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X}) Y_i}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{\sum_{i=1}^n (X_i - \bar{X})(\beta_0 + \beta_1 X_i + \epsilon_i)}{\sum_{i=1}^n (X_i - \bar{X})^2} = \beta_1 + \frac{\sum_{i=1}^n (X_i - \bar{X}) \epsilon_i}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

and so

$$\text{Var}(\hat{\beta}_1) = \text{Var}(\hat{\beta}_1 - \beta_1) = \frac{\sum_{i=1}^n (X_i - \bar{X})^2 \sigma^2}{(\sum_{i=1}^n (X_i - \bar{X})^2)^2} = \frac{\sigma^2}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

- (Optional) This immediately leads to an SE for $\hat{\beta}_1$ and shows that

$$\begin{aligned} E\left[\hat{\beta}_1^2 \sum_{i=1}^n (X_i - \bar{X})^2\right] &= E[\hat{\beta}_1^2] \sum_{i=1}^n (X_i - \bar{X})^2 \\ &= (\text{Var}(\hat{\beta}_1) + (E[\hat{\beta}_1])^2) \sum_{i=1}^n (X_i - \bar{X})^2 \\ &= \sigma^2 + \beta_1^2 \sum_{i=1}^n (X_i - \bar{X})^2 \end{aligned}$$

So, the Regression MS is also an estimate of σ^2 when $\beta_1 = 0$

Coefficient of Determination (or Multiple R-squared)

- Definition

$$R^2 = \frac{\text{Regression SS}}{\text{Total SS}} = 1 - \frac{\text{Residual SS}}{\text{Total SS}}$$

- Often expressed as a percentage
- A.k.a. Percentage of variance accounted for by regression
- Close to 1 (or 100%) if relationship is strong; precise prediction of individuals, as opposed to average, is possible
- For simple linear regression, R = coefficient of correlation
- For a given response variable, the smaller the Residual SS (i.e. the better the model fits) the larger R^2

Adjusted R^2 (or Adjusted R-squared)

- Definition

$$\text{Adjusted } R^2 = 1 - \frac{\text{Residual MS}}{\text{Total MS}} = 1 - \frac{\hat{\sigma}^2}{S_Y^2}$$

- For a given response variable, the smaller the estimate of σ^2 , the larger Adj R^2
- "Adjusted" for differing model df when there are many X variables and thus many models to choose from
- Penalizes R^2 for always increasing when more explanatory variables are put into the model
- It has no role in simple linear regression
- Adj R^2 can be negative

Simple Linear Regression

Towards a SE of a Fitted Value (Optional)

- We begin with our previous incarnation of $\hat{\beta}_1$, namely

$$\hat{\beta}_1 - \beta_1 = \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2} \sum_{i=1}^n (X_i - \bar{X}) \epsilon_i$$

and the variable \bar{Y} minus its expected value

$$\bar{Y} - \beta_0 - \beta_1 \bar{X} = \frac{1}{n} \sum_{i=1}^n \epsilon_i$$

Then

$$\begin{aligned} \text{Cov}(\bar{Y}, \hat{\beta}_1) &= E\left[(\bar{Y} - \beta_0 - \beta_1 \bar{X})(\hat{\beta}_1 - \beta_1)\right] \\ &= \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2} \frac{1}{n} \sum_{j=1}^n \sum_{i=1}^n (X_i - \bar{X}) E[\epsilon_i \epsilon_j] \\ &= \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2} \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X}) E[\epsilon_i^2] \\ &= \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2} \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X}) \sigma^2 = 0 \end{aligned}$$

Simple Linear Regression

SE of the Mean of Y at any X , of a Fitted Value and of β_0

- Knowing $Var(\bar{Y})$, $Var(\hat{\beta}_1)$ and $Cov(\bar{Y}, \hat{\beta}_1) = 0$ enables the calculation of the $Var(\hat{\mu}(Y|X_0))$ for any point $X = X_0$ (which includes $X =$ an observed X_i) as a simple sum of the variances:

$$\begin{aligned}Var[\hat{\mu}(Y|X_0)] &= Var[\hat{\beta}_0 + \hat{\beta}_1 X_0] = Var[\bar{Y} - \beta_1 \bar{X} + \hat{\beta}_1 X_0] \\&= Var[\bar{Y} - \hat{\beta}_1(\bar{X} - X_0)] \\&= \sigma^2 \left[\frac{1}{n} + \frac{(X_0 - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]\end{aligned}$$

- This leads to a SE for $\hat{\mu}(Y|X_0)$ which can be used to construct confidence intervals using the t_{n-2} distribution
- Putting $X = 0$ in the above expression gives

$$Var[\hat{\mu}(Y|X = 0)] = Var[\hat{\beta}_0] = \sigma^2 \left[\frac{1}{n} + \frac{\bar{X}^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]$$

Simple Linear Regression

Calculating Confidence Intervals for β_0 and β_1

- $100(1 - \alpha)\%$ Confidence intervals β_0 and β_1 are, respectively

$$\begin{aligned}\hat{\beta}_0 &\pm t_{n-2, \alpha/2} \times \sqrt{\hat{\sigma}^2 \left[\frac{1}{n} + \frac{\bar{X}^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]} \\ \hat{\beta}_1 &\pm t_{n-2, \alpha/2} \sqrt{\hat{\sigma}^2 \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2}}\end{aligned}$$

- In particular, we use the notation

$$\begin{aligned}\text{SE}(\hat{\beta}_0) &= \sqrt{\hat{\sigma}^2 \left[\frac{1}{n} + \frac{\bar{X}^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]} \\ \text{SE}(\hat{\beta}_1) &= \sqrt{\hat{\sigma}^2 \frac{1}{\sum_{i=1}^n (X_i - \bar{X})^2}}\end{aligned}$$

and $\hat{\sigma}^2$ is the Residual MS

Simple Linear Regression

Example: Ponds Institute

```
• > confint(F1, level=0.95)
```

		2.5 %	97.5 %
(Intercept)	38.3955349	49.02635	
price	-0.3135386	5.07000	

or

```
> ## need library("Rcmdr")  
> Confint(F1, level=0.95)
```

	Estimate		2.5 %	97.5 %
(Intercept)	43.71094	38.3955349	49.02635	
price	2.37823	-0.3135386	5.07000	

Simple Linear Regression

Calculating Confidence Intervals for the Mean of Y at Any Point X_0

- A $100(1 - \alpha)\%$ confidence interval for $\mu(Y|X_0)$ is

$$\hat{\beta}_0 + \hat{\beta}_1 X_0 \pm t_{n-2, \alpha/2} \sqrt{\hat{\sigma}^2 \left[\frac{1}{n} + \frac{(X_0 - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]}$$

- Here we use the notation

$$SE(\hat{\beta}_0 + \hat{\beta}_1 X_0) = \sqrt{\hat{\sigma}^2 \left[\frac{1}{n} + \frac{(X_0 - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]}$$

and $\hat{\sigma}^2$ is the Residual MS.

- In particular, when $X_0 = 0$,

$$SE(\hat{\beta}_0) = \sqrt{\hat{\sigma}^2 \left[\frac{1}{n} + \frac{\bar{X}^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]}$$

Simple Linear Regression

Example: Ponds Institute

- At $X_0 = 1$ and $X_0 = 3$:

```
> predict(F1,data.frame(price=c(1,3)),interval="confidence")
```

	fit	lwr	upr
1	46.08917	42.76392	49.41442
2	50.84563	46.35804	55.33322

Confidence Bands

- As X changes, the upper and lower bounds of the $100(1 - \alpha)\%$ C.I. for $\mu(Y|X)$ trace out curves in X

$$\hat{\beta}_0 + \hat{\beta}_1 X \pm t_{n-2, \alpha/2} \sqrt{\hat{\sigma}^2 \left[\frac{1}{n} + \frac{(X - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]}$$

These curves can be superimposed on the scatterplot of Y_i vs X_i to give some idea of the level of precision in the estimate of the $\mu(Y|X)$ at any point

Simple Linear Regression

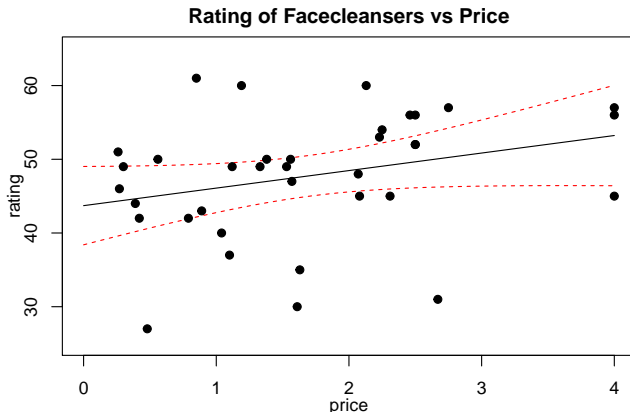
Example: Ponds Institute

```
> new = data.frame(price=seq(0,4,0.1))
> CIs = predict(F1,new,interval="confidence")
> matplot(new$price,CIs,lty=c(1,2,2),col=c("black","red","red"),
+         type="l",ylab="rating",main="Rating of Facecleansers vs Price",
+         ylim=c(25,65))
> points(rating~price,data=facecleanser,pch=21,bg="black")
```

Simple Linear Regression

Example: Ponds Institute

- The plot



Interval width increases with the distance from the mean of the X 's;
precision diminishes as the data becomes scarcer

Predicting The Impossible

- Up to this point we have been attempting to predict the average of Y at for any given X
- What if we were to predict a new observation at $X = X_0$?
- The new observation would be

$$Y_{\text{new}} = \beta_0 + \beta_1 X_0 + \epsilon$$

where ϵ is Normally distributed with mean 0 and variance σ^2 and independent of the sample values Y_i

- The estimate of Y_{new} , say

$$\hat{Y}_{\text{new}} = \hat{\beta}_0 + \hat{\beta}_1 X_0$$

Prediction Intervals

- the expected value and variance of the difference $Y_{\text{new}} - \hat{Y}_{\text{new}}$

$$\begin{aligned}E(Y_{\text{new}} - \hat{Y}_{\text{new}}) &= 0 \\ \text{Var}(Y_{\text{new}} - \hat{Y}_{\text{new}}) &= \sigma^2 + \text{Var}(\hat{Y}_{\text{new}})\end{aligned}$$

where

$$\text{Var}(\hat{Y}_{\text{new}}) = \sigma^2 \left[\frac{1}{n} + \frac{(X_0 - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]$$

- Estimating σ^2 by the Residual MS leads to a t_{n-2} distribution for $Y_{\text{new}} - \hat{Y}_{\text{new}}$, which gives the $100(1 - \alpha)\%$ CI for Y_{new} as

$$\hat{\beta}_0 + \hat{\beta}_1 X_0 \pm t_{n-2, \alpha/2} \sqrt{\hat{\sigma}^2 \left[1 + \frac{1}{n} + \frac{(X_0 - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right]}$$

Such an interval is usually termed a $100(1 - \alpha)\%$ prediction interval, to distinguish it from a confidence interval of a fixed parameter

Prediction Intervals

- As $\hat{\sigma}^2 \left[\frac{1}{n} + \frac{(X_0 - \bar{X})^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right] \rightarrow 0$ when $n \rightarrow \infty$, The first term dominates the SE and the prediction interval is effectively

$$\hat{\beta}_0 + \hat{\beta}_1 X_0 \pm t_{n-2, \alpha/2} \sqrt{\hat{\sigma}^2}$$

as might be expected from the assumption that the variance of an observation at any given X is constant

R^2 Or Not R^2

- For some regressionistas, R^2 dominates the discussion of results and no model will be countenanced unless it exceeds 90%
- Such dizzy heights are seldom seen in any data collected from living things
- When considering summary output, you need to bear in mind the aim of the regression
- Certainly R^2 needs to be very high if you're contemplating estimating a new observation
- For those who merely wish to establish a relationship or are happy estimating the mean response, R^2 can take a back seat
- Also be mindful that R^2 is on a squared scale, so that an R^2 of 50% corresponds to an r of 0.7, which is a relatively high correlation; almost certainly significant

Simple Linear Regression

Example: Cholesterol Data Set

- Data from 1109 West Australians with measurements pertaining to body mass, cholesterol, blood pressure and other data of medical obsession
- Is BMI related to cholesterol?

```
> load("cholesterol.RData")
> str(cholesterol)

'data.frame': 1109 obs. of  8 variables:
 $ AGE      : num  32 40 39 37 46 44 51 50 49 49 ...
 $ BMI      : num  24.2 26.3 25.1 28.7 26.3 ...
 $ CHOL     : num  4.7 5.8 5.5 5.6 5.9 5.8 5.4 6 4.9 7.2 ...
 $ DBP      : num  70 70 70 80 80 84 90 80 84 90 ...
 $ HEIGHT   : num  175 183 182 183 185 180 170 173 187 178 ...
 $ SEX      : Factor w/ 2 levels "female","male": 2 2 2 2 2 2 2 2 2 2 ...
 $ WAIST    : num  82 93 91 98 95 84 82 92 95 89 ...
 $ WEIGHT   : num  74 88 83 96 90 76 72 75 98 89 ...
 - attr(*, "variable.labels")= Named chr  "age" "BMI" "cholesterol" "DBP" ...
 ..- attr(*, "names")= chr  "AGE" "BMI" "CHOL" "DBP" ...
```

Simple Linear Regression

Example: Cholesterol Data Set


```
• > cholesterol[1:18,]
```

	AGE	BMI	CHOL	DBP	HEIGHT	SEX	WAIST	WEIGHT
1	32	24.16	4.7	70	175	male	82	74
2	40	26.28	5.8	70	183	male	93	88
3	39	25.06	5.5	70	182	male	91	83
4	37	28.67	5.6	80	183	male	98	96
5	46	26.30	5.9	80	185	male	95	90
6	44	23.46	5.8	84	180	male	84	76
7	51	24.91	5.4	90	170	male	82	72
8	50	25.06	6.0	80	173	male	92	75
9	49	28.02	4.9	84	187	male	95	98
10	49	28.09	7.2	90	178	male	89	89
11	55	26.51	5.1	100	178	male	105	84
12	52	29.98	4.8	80	178	male	102	95
13	54	29.32	7.1	80	180	male	107	95
14	61	31.26	4.6	80	185	male	109	107
15	59	24.22	6.4	70	170	male	92	70
16	61	23.77	4.7	90	180	male	89	77
17	67	28.73	6.3	80	175	male	98	88
18	67	25.03	8.3	90	181	male	97	82

Simple Linear Regression

Example: Cholesterol Data Set

- The plot



This is a large data set of high variability. The sample size works in favour of the CI (red), but the PI (blue) is left out in the cold.

Simple Linear Regression

Example: Cholesterol Data Set

- The code

```
> C1 = lm(BMI~CHOL,data=cholesterol)
> plot(BMI~CHOL,data=cholesterol,pch=21,bg="black",main="BMI vs Cholesterol")
> new = data.frame(CHOL=seq(2,10,0.1))
> CIs = predict(C1,new,interval="confidence")
> PIs = predict(C1,new,interval="predict")
> matpoints(new$CHOL,CIs,lty=c(1,2,2),col=c("black","red","red"),type="l")
> matpoints(new$CHOL,PIs,lty=c(1,2,2),col=c("black","blue","blue"),type="l")
```

Simple Linear Regression

Example: Cholesterol Data Set

• The Fitted Model

```
> C1 = lm(BMI~CHOL,data=cholesterol)
> anova(C1) ## Analysis of Variance Table
```

Analysis of Variance Table

Response: BMI

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
CHOL	1	184.3	184.333	11.141	0.0008726 ***
Residuals	1107	18316.1	16.546		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Simple Linear Regression

Example: Cholesterol Data Set

- The Fitted Model

```
> summary(C1)
```

Call:

```
lm(formula = BMI ~ CHOL, data = cholesterol)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.5632	-2.9024	-0.4118	2.4582	15.3019

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	23.8078	0.6919	34.409	< 2e-16 ***
CHOL	0.4086	0.1224	3.338	0.000873 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.068 on 1107 degrees of freedom

Multiple R-squared: 0.009964, Adjusted R-squared: 0.009069

F-statistic: 11.14 on 1 and 1107 DF, p-value: 0.0008726

Simple Linear Regression

Example: Cholesterol Data Set

- Conclusion

- ```
> confint(C1,level=0.95)
```

|             | 2.5 %      | 97.5 %     |
|-------------|------------|------------|
| (Intercept) | 22.4501727 | 25.1653519 |
| CHOL        | 0.1683972  | 0.6487605  |

or

```
> ## need library("Rcmdr")
> Confint(C1,level=0.95)
```

|             | Estimate   | 2.5 %      | 97.5 %     |
|-------------|------------|------------|------------|
| (Intercept) | 23.8077623 | 22.4501727 | 25.1653519 |
| CHOL        | 0.4085789  | 0.1683972  | 0.6487605  |

As cholesterol increases by 1, mean BMI is estimated to increase by between 0.168 and 0.649

- Note there is no causal relationship implied or intended to be implied by this statement



## Example: Cholesterol Data Set

- Conclusion
  - Whether such an increase is of practical significance is not something statistics has an opinion on
  - Clearly in this example, there is a great deal of "unexplained" variability; 99% in fact
  - There are many other variables involved in determining BMI. Incorporating those variables into the equation to obtain a better explanation is the province of multiple regression
- Disclaimer
  - The interpretations of the examples presented so far are only valid provided the model assumptions are valid
  - Nothing in the numbers presented so far will tell you if this is the case or not

## Anscombe Quartet


```
> load("anscombe.RData")
> anscombe
```

|    | x1 | y1    | x2 | y2   | x3 | y3    | x4 | y4    |
|----|----|-------|----|------|----|-------|----|-------|
| 1  | 4  | 4.26  | 4  | 3.10 | 4  | 5.39  | 8  | 5.56  |
| 2  | 5  | 5.68  | 5  | 4.74 | 5  | 5.73  | 19 | 12.50 |
| 3  | 6  | 7.24  | 6  | 6.13 | 6  | 6.08  | 8  | 5.25  |
| 4  | 7  | 4.82  | 7  | 7.26 | 7  | 6.42  | 8  | 6.89  |
| 5  | 8  | 6.95  | 8  | 8.14 | 8  | 6.77  | 8  | 5.76  |
| 6  | 9  | 8.81  | 9  | 8.77 | 9  | 7.11  | 8  | 8.84  |
| 7  | 10 | 8.04  | 10 | 9.14 | 10 | 7.46  | 8  | 6.58  |
| 8  | 11 | 8.33  | 11 | 9.26 | 11 | 7.81  | 8  | 8.47  |
| 9  | 12 | 10.84 | 12 | 9.13 | 12 | 8.15  | 8  | 7.91  |
| 10 | 13 | 7.58  | 13 | 8.74 | 13 | 12.74 | 8  | 7.71  |
| 11 | 14 | 9.96  | 14 | 8.10 | 14 | 8.84  | 8  | 7.04  |

# Simple Linear Regression

## Anscombe Quartet


- Scatterplots



# Simple Linear Regression

## Anscombe Quartet

- The fits



Equation of line is  $y = 3 + 0.5x$  with  $R^2 = 67\%$  in each case

## Virtually Identical Output

### • R-output

```
> A1 = lm(y1~x1,data=anscombe)
> summary(A1)
```

Call:

```
lm(formula = y1 ~ x1, data = anscombe)
```

Residuals:

| Min      | 1Q       | Median   | 3Q      | Max     |
|----------|----------|----------|---------|---------|
| -1.92127 | -0.45577 | -0.04136 | 0.70941 | 1.83882 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 3.0001   | 1.1247     | 2.667   | 0.02573 *  |
| x1          | 0.5001   | 0.1179     | 4.241   | 0.00217 ** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.237 on 9 degrees of freedom

Multiple R-squared: 0.6665, Adjusted R-squared: 0.6295

F-statistic: 17.99 on 1 and 9 DF, p-value: 0.00217

## Virtually Identical Output

### • R-output

```
> A2 = lm(y2~x2,data=anscombe)
> summary(A2)
```

Call:

```
lm(formula = y2 ~ x2, data = anscombe)
```

Residuals:

| Min     | 1Q      | Median | 3Q     | Max    |
|---------|---------|--------|--------|--------|
| -1.9009 | -0.7609 | 0.1291 | 0.9491 | 1.2691 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 3.001    | 1.125      | 2.667   | 0.02576 *  |
| x2          | 0.500    | 0.118      | 4.239   | 0.00218 ** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.237 on 9 degrees of freedom

Multiple R-squared: 0.6662, Adjusted R-squared: 0.6292

F-statistic: 17.97 on 1 and 9 DF, p-value: 0.002179

## Virtually Identical Output

### • R-output

```
> A3 = lm(y3~x3,data=anscombe)
> summary(A3)
```

Call:

```
lm(formula = y3 ~ x3, data = anscombe)
```

Residuals:

| Min     | 1Q      | Median  | 3Q     | Max    |
|---------|---------|---------|--------|--------|
| -1.1586 | -0.6146 | -0.2303 | 0.1540 | 3.2411 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 3.0025   | 1.1245     | 2.670   | 0.02562 *  |
| x3          | 0.4997   | 0.1179     | 4.239   | 0.00218 ** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.236 on 9 degrees of freedom

Multiple R-squared: 0.6663, Adjusted R-squared: 0.6292

F-statistic: 17.97 on 1 and 9 DF, p-value: 0.002176

## Virtually Identical Output

- R-output

```
> A4 = lm(y4~x4,data=anscombe)
> summary(A4)
```

Call:

```
lm(formula = y4 ~ x4, data = anscombe)
```

Residuals:

| Min    | 1Q     | Median | 3Q    | Max   |
|--------|--------|--------|-------|-------|
| -1.751 | -0.831 | 0.000  | 0.809 | 1.839 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )   |
|-------------|----------|------------|---------|------------|
| (Intercept) | 3.0017   | 1.1239     | 2.671   | 0.02559 *  |
| x4          | 0.4999   | 0.1178     | 4.243   | 0.00216 ** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.236 on 9 degrees of freedom

Multiple R-squared: 0.6667, Adjusted R-squared: 0.6297

F-statistic: 18 on 1 and 9 DF, p-value: 0.002165



## Important Lesson

- Nothing in the regression output, be it  $R^2$ , adjusted  $R^2$ , Residual MS,  $F$  ratio,  $p$ -value, regression coefficients or any other number, will alert you to possibly serious inadequacies in the analysis
- The adequacy of the fitted model and potential problems must be assessed independently of standard output
- Diagnostics:
  - Plot the data
  - Plot Standardized Residuals vs Fitted Values
  - Plot Standardized Residuals vs  $X$
  - Leverages and Cook's Distances
  - Q-Q plot of Standardized Residuals

### Plotting Standardized Residuals vs Fitted Values or $X$


- Check Assumptions and Model Fit:
  - Check constant variance; do the points fan out as  $X$  increases
  - Check for large residuals; 95% should lie within the tramlines of  $y = \pm 2$
  - Check for skewness and model fit; equal numbers of positive and negative residuals at each  $X$
  - Check the adequacy of the straight line to model the data; there should be no curvature apparent in the plot
  - Check for unusual patterns or clusters; the plot should look as "random" as a shotgun blast
  - Check for points of high leverage; points whose  $X$  value is extreme

# Simple Linear Regression

Example: Cholesterol Data Set

- BMI on CHOL Standardized Residuals vs Fits Plot

```
> plot(rstandard(C1)~fitted(C1),pch=21,bg="black",main="Standardized Residuals
+ xlab="Fitted Values",ylab="Standardized Residuals")
> abline(h=c(-2,0,2),col=c(2,1,2),lty=2)
```



Many overly large positive residuals. Skew distribution of standardized residuals. Some evidence of increasing variance.


# Simple Linear Regression

Example: Cholesterol Data Set


- BMI on CHOL QQ Plot of Standardized Residual

```
> hist(rstandard(C1))
> qqnorm(rstandard(C1))
> qqline(rstandard(C1))
```

Histogram of rstandard(C



Normal Q-Q Plot



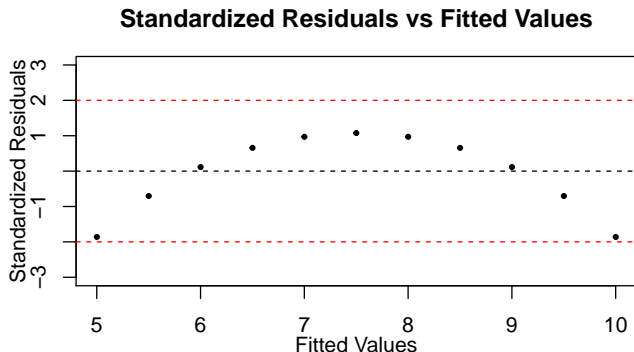
Residuals are skewed to right; not Normal. A log transform of BMI and possibly also CHOL is suggested.

# Simple Linear Regression

## Example

- Anscombe's Y2 on X2 Standardized Residuals vs Fits Plot

```
> plot(rstandard(A2)~fitted(A2),pch=21,bg="black",main="Standardized Residuals
+ xlab="Fitted Values",ylab="Standardized Residuals")
> abline(h=c(-2,0,2),col=c(2,1,2),lty=2)
```




Quadratic trend in the residuals. Straight line fit is inadequate. Add quadratic term to the model.

# Simple Linear Regression

## Example

- Anscombe's Y3 on X3 Standardized Residuals vs Fits Plot

```
> plot(rstandard(A3)~fitted(A3),pch=21,bg="black",main="Standardized Residuals
+ xlab="Fitted Values",ylab="Standardized Residuals")
> abline(h=c(-2,0,2),col=c(2,1,2),lty=2)
```



Outlier has pulled line towards it. Linear trend still exists in the residuals. Fitted slope does not match the majority of the data.


Problematic Points of Two Types (outliers or/and points of high leverage)

- Individual points in a regression can be problematic because their  $Y$  value is a long way from where it should be (such points are outliers)
- Equally problematic are points whose  $X$  value is a long way from the bulk of the other  $X$ 's (these are points of high leverage)
- Presence of either type of points can have a substantial effect on the analysis and its conclusions
- Often their presence indicates that an inadequate model has been fitted and should not immediately set in train a massive cull of offending points

# Simple Linear Regression

Scatterplots of Standardized  $X$  and  $Y$  with Various Values of  $r_{XY}$

- High Leverage Points



Points at some remove from the other  $X$ 's. (Whither thou goest I will go)



## Towards the Leverages

- The leverage of a sample point  $(X_i, Y_i)$  is derived from the variance of its residual

$$\text{Var}(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = \sigma^2 + \text{Var}(\hat{\beta}_0 + \hat{\beta}_1 X_i) - 2\text{Cov}(Y_i, \hat{\beta}_0 + \hat{\beta}_1 X_i)$$

### Towards the Leverages (Optional)

- The variance term we know as the variance of a fitted value

$$\text{Var}(\hat{\beta}_0 + \hat{\beta}_1 X_i) = \sigma^2 \left[ \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum_{j=1}^n (X_j - \bar{X})^2} \right]$$

- The covariance term involves the product of

$$\begin{aligned} Y_i - E[Y_i] &= \epsilon_i \\ \hat{\beta}_0 + \hat{\beta}_1 X_i - E[\hat{\beta}_0 + \hat{\beta}_1 X_i] &= (\hat{\beta}_0 - \beta_0) + (\hat{\beta}_1 - \beta_1) X_i \\ &= \frac{1}{n} \sum_{j=1}^n \epsilon_j + (\hat{\beta}_1 - \beta_1)(X_i - \bar{X}) \end{aligned}$$

because

$$\begin{aligned} (\hat{\beta}_0 + \hat{\beta}_1 \bar{X}) - \bar{Y} &= (\hat{\beta}_0 + \hat{\beta}_1 \bar{X}) - (\beta_0 + \beta_1 \bar{X} + \frac{1}{n} \sum_{j=1}^n \epsilon_j) \\ \Rightarrow \hat{\beta}_0 - \beta_0 &= \frac{1}{n} \sum_{j=1}^n \epsilon_j - (\hat{\beta}_1 - \beta_1) \bar{X} \end{aligned}$$

## Towards the Leverages (Optional)

- The calculation uses the " $\epsilon$ " formulation of

$$\hat{\beta}_1 - \beta_1 = \frac{\sum_{j=1}^n (X_j - \bar{X}) \epsilon_j}{\sum_{j=1}^n (X_j - \bar{X})^2}$$

- The Covariance Term

$$\begin{aligned} \text{Cov}(Y_i, \hat{\beta}_0 + \hat{\beta}_1 X_i) &= E\left[(Y_i - E[Y_i])(\hat{\beta}_0 + \hat{\beta}_1 X_i - E[\hat{\beta}_0 + \hat{\beta}_1 X_i])\right] \\ &= E\left[\epsilon_i \left(\frac{1}{n} \sum_{j=1}^n \epsilon_j + (\hat{\beta}_1 - \beta_1)(X_i - \bar{X})\right)\right] \\ &= E\left[\epsilon_i \left(\frac{1}{n} \sum_{j=1}^n \epsilon_j + \frac{\sum_{j=1}^n (X_j - \bar{X}) \epsilon_j}{\sum_{j=1}^n (X_j - \bar{X})^2} (X_i - \bar{X})\right)\right] \\ &= \frac{1}{n} \sum_{j=1}^n E[\epsilon_i \epsilon_j] + \frac{\sum_{j=1}^n (X_j - \bar{X})(X_i - \bar{X}) E[\epsilon_i \epsilon_j]}{\sum_{j=1}^n (X_j - \bar{X})^2} \\ &= \frac{1}{n} E[\epsilon_i^2] + \frac{(X_i - \bar{X})^2 E[\epsilon_i^2]}{\sum_{j=1}^n (X_j - \bar{X})^2} \\ &= \sigma^2 \left[ \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum_{j=1}^n (X_j - \bar{X})^2} \right] \end{aligned}$$

### Variance of a Residual and Expected Residual MS

- The variance of the residual at  $(X_i, Y_i)$  is thus

$$\text{Var}(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = \left[ 1 - \left( \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum_{j=1}^n (X_j - \bar{X})^2} \right) \right] \sigma^2$$

- (Optional) This not only gives us the leverages but also clears an item in the pending tray:

$$\begin{aligned} E[\text{Residual SS}] &= E\left[\sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i)^2\right] \\ &= \sum_{i=1}^n E\left[(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i - E[Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i])^2\right] \\ &= \sum_{i=1}^n \text{Var}(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) \\ &= \left[ n - \frac{n}{n} - \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{\sum_{j=1}^n (X_j - \bar{X})^2} \right] \sigma^2 = (n - 2) \sigma^2 \end{aligned}$$

showing the Residual MS is an unbiased estimate of  $\sigma^2$

# Simple Linear Regression

## Leverages

- The  $i$ th leverage is denoted  $h_i$ . By definition

$$\text{Var}(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i) = (1 - h_i)\sigma^2$$

and so

$$h_i = \frac{1}{n} + \frac{(X_i - \bar{X})^2}{\sum_{j=1}^n (X_j - \bar{X})^2}$$

- Some maximization under constraint shows  $0 < h_i < 1$
- The further  $X_i$  is from  $\bar{X}$  the closer its leverage to 1
- If the leverage is very close to 1, the residual is 0 with probability close to 1 and the point's predicted value is close to its observed value, irrespective of any input from the rest of the data
- A point is defined to have high leverage if

$$h_i > 2 \times \text{average leverage} = 2\frac{p}{n}$$

where  $p = \text{Regression df} + 1 = 2$  for simple linear regression

## Example: Huber's Data

- `> load("huber.RData")`

- `> huber`

|   | x  | BadY  | GoodY  |
|---|----|-------|--------|
| 1 | -4 | 2.48  | 2.48   |
| 2 | -3 | 0.73  | 0.73   |
| 3 | -2 | -0.04 | -0.04  |
| 4 | -1 | -1.44 | -1.44  |
| 5 | 0  | -1.32 | -1.32  |
| 6 | 10 | 0.00  | -11.40 |

- The point of contention is at  $X = 10$ ; otherwise  $\text{GoodY} = \text{BadY}$

# Simple Linear Regression

## Example: Huber's Data

- $n = 6$  and so the leverage cut-off point is  $4/6 = 0.667$

```
> H1 = lm(GoodY~x,data=huber)
> H2 = lm(BadY~x,data=huber)
> cbind('H1 stdres'=rstandard(H1), 'H1 hats'=hatvalues(H1),
+ 'H2 stdres'=rstandard(H2), 'H2 hats'=hatvalues(H2))
```


|   | H1 stdres  | H1 hats   | H2 stdres  | H2 hats   |
|---|------------|-----------|------------|-----------|
| 1 | 1.1792919  | 0.2897436 | 1.5967356  | 0.2897436 |
| 2 | -0.7454671 | 0.2358974 | 0.3079765  | 0.2358974 |
| 3 | -0.2902758 | 0.1974359 | -0.1953457 | 0.1974359 |
| 4 | -1.2964502 | 0.1743590 | -1.1287795 | 0.1743590 |
| 5 | 1.1650975  | 0.1666667 | -0.9811785 | 0.1666667 |
| 6 | 0.1273613  | 0.9358974 | 1.9016428  | 0.9358974 |

leverages (or hat values) are the same because the  $X$ 's are the same


# Simple Linear Regression

## Example: Huber's Data

- Comparison (for the 6th point)



$\text{stdres} = 0.127$  &  $\text{hats} = 0.936$   
Point follows the trend of the rest of the data  
It's a point of high leverage but not an outlier  
No action necessary



$\text{stdres} = 1.902$  &  $\text{hats} = 0.936$   
Point does not follow the trend of the rest of the data  
It's a point of high leverage and an outlier ( $|\text{stdres}| > 2$ , almost)  
Fit a different model




# Simple Linear Regression


Example: Huber's Data

- Comparison (Better Pictures)

Model: H1



Model: H2



# Simple Linear Regression

## Example: Huber's Data

### • Code

```
> plot(rstandard(H1)~hatvalues(H1),pch=16,xlim=c(0,1),ylim=c(-3,3),
+ xlab="Hat values",ylab="Standard Residuals")
> abline(h=c(-2,2),col=c(2,2),lty=2)
> abline(v=4/6,col=2,lty=2)
> plot(rstandard(H2)~hatvalues(H2),pch=16,xlim=c(0,1),ylim=c(-3,3),
+ xlab="Hat values",ylab="Standard Residuals")
> abline(h=c(-2,2),col=c(2,2),lty=2)
> abline(v=4/6,col=2,lty=2)
```

### Huber's Process and Standardization

- A point can have high leverage or not, or have a standardized residual exceeding 2 (roughly) in magnitude or not
- From a leverage perspective, only high leverage points which are also outliers need actioning
- This is not to say outliers per se don't need actioning
- It is important that the standardized residuals used in determining whether a point of high leverage is good or bad are the individually standardized residuals and not globally standardized residuals. E.g. Excel and SPSS both use  $\sqrt{\hat{\sigma}^2}$  and not  $\sqrt{(1 - h_i)\hat{\sigma}^2}$  when "standardizing". SPSS and The Statistical Sleuth reserve the terminology "studentized" for their properly standardized residuals

## Simple Linear Regression

### Cook's Distance – Combining Leverage & Standardized Residual Information

- The good leverage point, bad leverage point process works well except when there is no relationship among the rest of data. E.g. In Anscombe's 4th data set the lone point has residual = 0 and leverage = 1

```
> cbind('y4'=anscombe$y4, 'x4'=anscombe$x4,
+ 'A4 res'=residuals(A4), 'A4 hats'=hatvalues(A4))
```

|    | y4    | x4 | A4 res | A4 hats |
|----|-------|----|--------|---------|
| 1  | 5.56  | 8  | -1.441 | 0.1     |
| 2  | 12.50 | 19 | 0.000  | 1.0     |
| 3  | 5.25  | 8  | -1.751 | 0.1     |
| 4  | 6.89  | 8  | -0.111 | 0.1     |
| 5  | 5.76  | 8  | -1.241 | 0.1     |
| 6  | 8.84  | 8  | 1.839  | 0.1     |
| 7  | 6.58  | 8  | -0.421 | 0.1     |
| 8  | 8.47  | 8  | 1.469  | 0.1     |
| 9  | 7.91  | 8  | 0.909  | 0.1     |
| 10 | 7.71  | 8  | 0.709  | 0.1     |
| 11 | 7.04  | 8  | 0.039  | 0.1     |

### Cook's Distance – Combining Leverage & Standardized Residual Information

- For a data point  $(X_i, Y_i)$ , Cook's Distance

$$D_i = p^{-1}(\text{standardized residual}_i)^2 \times \frac{h_i}{1 - h_i}$$

where  $p = \text{Regression df} + 1 = 2$  in simple linear regression

- Cook's Distance is large if a point has a large standardized residual, a large leverage or both
- If  $D_i > 1$  the point is considered to have undue influence
- Calculate the predicted values from the regression including case  $i$  and the predicted values from the regression excluding case  $i$
- Calculate the sum of the squared differences between these predictions and divide by  $p\hat{\sigma}^2$
- This is  $D_i$ , Cook's distance for case  $i$

# Simple Linear Regression

## Example: Huber's Data (Cook's Distances in R)

```
> cooks.distance(H1)
 1 2 3 4 5 6
0.28366861 0.08578246 0.01036426 0.17747401 0.13574522 0.11841263
> (1/2)*rstandard(H1)^2*hatvalues(H1)/(1-hatvalues(H1))
 1 2 3 4 5 6
0.28366861 0.08578246 0.01036426 0.17747401 0.13574522 0.11841263
> cooks.distance(H2)
 1 2 3 4 5 6
0.520037540 0.014641199 0.004693794 0.134536856 0.096271128 26.398591892
> (1/2)*rstandard(H2)^2*hatvalues(H2)/(1-hatvalues(H2))
 1 2 3 4 5 6
0.520037540 0.014641199 0.004693794 0.134536856 0.096271128 26.398591892
```

The not quite bad but definitely naughty leverage point of model H2 is clearly overstepping Cook's line in the sand

## Simple Linear Regression

### Example: Huber's Data (Alternative Derivation of Cook's Distance)

- Alternatively you can get a raft of influence measures, together with a flag if any overstep some line

```
> influence.measures(H2)
```

```
Influence measures of
```

```
lm(formula = BadY ~ x, data = huber) :
```

|   | dfb.1_  | dfb.x     | dffit   | cov.r | cook.d   | hat   | inf |
|---|---------|-----------|---------|-------|----------|-------|-----|
| 1 | 1.1124  | -9.56e-01 | 1.4667  | 0.329 | 0.52004  | 0.290 | *   |
| 2 | 0.1261  | -8.13e-02 | 0.1500  | 2.218 | 0.01464  | 0.236 |     |
| 3 | -0.0775 | 3.33e-02  | -0.0843 | 2.173 | 0.00469  | 0.197 |     |
| 4 | -0.5320 | 1.14e-01  | -0.5442 | 1.000 | 0.13454  | 0.174 |     |
| 5 | -0.4361 | 1.78e-17  | -0.4361 | 1.230 | 0.09627  | 0.167 |     |
| 6 | 8.5733  | 1.84e+01  | 20.3160 | 0.255 | 26.39859 | 0.936 | *   |

# Simple Linear Regression

## Example: Anscombe's Y4 on X4

```
• > influence.measures(A4)
```

Influence measures of

lm(formula = y4 ~ x4, data = anscombe) :

|    | dfb.1_   | dfb.x4   | dffit   | cov.r | cook.d   | hat | inf |
|----|----------|----------|---------|-------|----------|-----|-----|
| 1  | -0.25432 | 0.12769  | -0.4235 | 0.974 | 8.39e-02 | 0.1 |     |
| 2  | 0.00000  | 0.00000  | NaN     | NaN   | NaN      | 1.0 | *   |
| 3  | -0.32505 | 0.16320  | -0.5413 | 0.795 | 1.24e-01 | 0.1 |     |
| 4  | -0.01788 | 0.00898  | -0.0298 | 1.403 | 4.98e-04 | 0.1 |     |
| 5  | -0.21353 | 0.10721  | -0.3556 | 1.078 | 6.23e-02 | 0.1 |     |
| 6  | 0.34734  | -0.17439 | 0.5784  | 0.742 | 1.37e-01 | 0.1 |     |
| 7  | -0.06827 | 0.03428  | -0.1137 | 1.366 | 7.17e-03 | 0.1 |     |
| 8  | 0.26029  | -0.13069 | 0.4334  | 0.958 | 8.72e-02 | 0.1 |     |
| 9  | 0.15149  | -0.07606 | 0.2523  | 1.225 | 3.34e-02 | 0.1 |     |
| 10 | 0.11654  | -0.05851 | 0.1941  | 1.294 | 2.03e-02 | 0.1 |     |
| 11 | 0.00628  | -0.00315 | 0.0105  | 1.406 | 6.15e-05 | 0.1 |     |



# Simple Linear Regression

## Example: Cargo Data

- Relationship between the volume of a ship's cargo loaded and unloaded ( $X$ ) and the time in hours spent in port ( $Y$ )

```
> load("cargoes.RData")
> str(cargoes)

'data.frame': 31 obs. of 2 variables:
 $ Tonnage: num 268 294 329 353 363 507 529 536 547 663 ...
 $ Time : num 11 13 13 15 20 11 11 22 20 13 ...
```

# Simple Linear Regression

## Example: Cargo Data

```
• > cargoes[1:18,]
```

|    | Tonnage | Time |
|----|---------|------|
| 1  | 268     | 11   |
| 2  | 294     | 13   |
| 3  | 329     | 13   |
| 4  | 353     | 15   |
| 5  | 363     | 20   |
| 6  | 507     | 11   |
| 7  | 529     | 11   |
| 8  | 536     | 22   |
| 9  | 547     | 20   |
| 10 | 663     | 13   |
| 11 | 851     | 9    |
| 12 | 1328    | 15   |
| 13 | 1486    | 28   |
| 14 | 1732    | 24   |
| 15 | 1849    | 17   |
| 16 | 2213    | 17   |
| 17 | 2790    | 43   |
| 18 | 2829    | 30   |

# Simple Linear Regression

## Example: Cargo Data


```
• > cargoes[19:31,]
```

|    | Tonnage | Time |
|----|---------|------|
| 19 | 3192    | 55   |
| 20 | 3256    | 30   |
| 21 | 3930    | 43   |
| 22 | 4263    | 31   |
| 23 | 4682    | 49   |
| 24 | 5375    | 68   |
| 25 | 6112    | 69   |
| 26 | 6666    | 49   |
| 27 | 6760    | 43   |
| 28 | 7021    | 64   |
| 29 | 7084    | 41   |
| 30 | 12203   | 68   |
| 31 | 15900   | 131  |

# Simple Linear Regression

## Example: Cargo Data

```
> with(cargoes,plot(Time~Tonnage,xlim=c(0,16000),ylim=c(0,140),pch=16))
> with(cargoes,lines(fitted(lm(Time~Tonnage))~Tonnage,col=2))
```



Variance appears to be increasing and there are two points of obvious high leverage - but will they be good or bad?

# Simple Linear Regression

## Example: Cargo Data

```
> D1 = lm(Time~Tonnage,data=cargoes)
> summary(D1)
```

Call:

```
lm(formula = Time ~ Tonnage, data = cargoes)
```

Residuals:

| Min     | 1Q     | Median | 3Q    | Max    |
|---------|--------|--------|-------|--------|
| -23.882 | -6.397 | -1.261 | 5.931 | 21.850 |

Coefficients:

|             | Estimate  | Std. Error | t value | Pr(> t ) |     |
|-------------|-----------|------------|---------|----------|-----|
| (Intercept) | 12.344707 | 2.642633   | 4.671   | 6.32e-05 | *** |
| Tonnage     | 0.006518  | 0.000531   | 12.275  | 5.22e-13 | *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.7 on 29 degrees of freedom

Multiple R-squared: 0.8386, Adjusted R-squared: 0.833


F-statistic: 150.7 on 1 and 29 DF, p-value: 5.218e-13

What part of this output can be trusted, apart from the arithmetic?

# Simple Linear Regression

## Example: Cargo Data

```
> plot(rstandard(D1)~fitted(D1),pch=21,bg="black",main="Standardized Residuals vs Fitted Values",
+ xlab="Fitted Values",ylab="Standardized Residuals")
> abline(h=c(-2,0,2),col=c(2,1,2),lty=2)
```




A clear case of increasing variance, with both high leverage points living dangerously. No obvious evidence of curvature.

# Simple Linear Regression

## Example: Cargo Data

```
> plot(hatvalues(D1)~c(1:31),pch=21,bg="black",main="Leverages vs Row Numbers",
+ xlab="Row Numbers",ylab="Leverages")
> abline(h=2*2/31,col=2,lty=2)
```

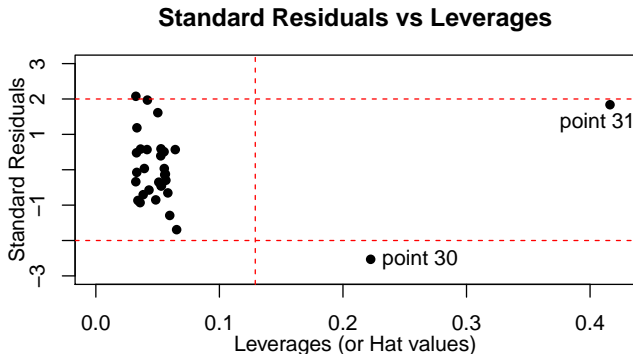


The line in the sand is  $2 \times 2/31 = 0.129$ .

# Simple Linear Regression

## Example: Cargo Data

```
> plot(rstandard(D1)~hatvalues(D1),pch=16,main="Standard Residuals vs Leverages",
+ xlab="Leverages",ylab="Standard Residuals")
> abline(h=c(-2,2),col=c(2,2),lty=2)
> abline(v=2*2/31,col=2,lty=2)
```




Point 30 is definitely bad leverage point and 31 somewhat naughty to say the least.



# Simple Linear Regression

## Example: Cargo Data

```
> plot(cooks.distance(D1)~c(1:31),pch=21,bg="black",main="Cooks' distance vs Row Numbers",
+ xlab="Row Numbers",ylab="Cooks' distance")
> abline(h=1,col=2,lty=2)
```




Point 31 is now in influential and 30 just under the radar. Both probably need a stern talking to.

# Simple Linear Regression

## Example: Cargo Data

- One from the `plot(D1)` Stable

```
> plot(D1,5,add.smooth=FALSE,pch=16)
```




A combination standardized residual, leverage and Cook's distance plot. Points beyond the red lines are influential. (better to use `plot(D1)` to get all the plots instead)

# Simple Linear Regression

## Example: Cargo Data

- The plot



Variance is overestimated at low tonnages and underestimated at high tonnages. The latter is worse in the context. (What Can Be Done? Transformation)

# Simple Linear Regression

## Example: Cargo Data

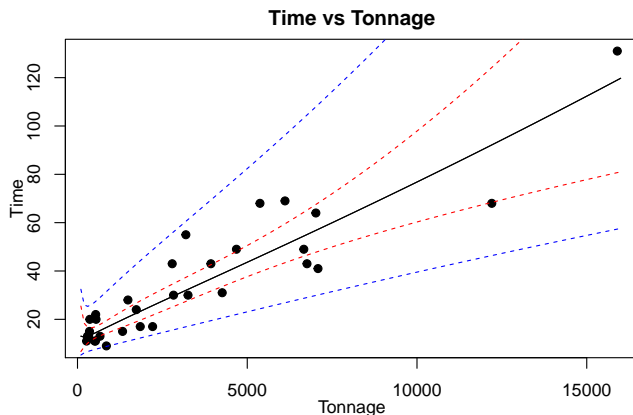
- The code

```
> plot(Time~Tonnage,data=cargoes,pch=21,bg="black",main="Time vs Tonnage")
> new = data.frame(Tonnage=seq(0,16000,100))
> CIs = predict(D1,new,interval="confidence")
> PIs = predict(D1,new,interval="predict")
> matpoints(new$Tonnage,CIs,lty=c(1,2,2),col=c("black","red","red"),type="l")
> matpoints(new$Tonnage,PIs,lty=c(1,2,2),col=c("black","blue","blue"),type="l")
```

# Simple Linear Regression

## Example: Cargo Data

- log transformation in both sides



Quadratic fit of  $\log(\text{Time})$  on  $\log(\text{Tonnage})$  transformed back to the original scale. The fitted line is almost identical but the estimates of precision very different. (need to check assumptions for sure)

# Simple Linear Regression

Example:

- The code

```
> plot(Time~Tonnage,data=cargoes,pch=21,bg="black",main="Time vs Tonnage")
> new = data.frame(Tonnage=seq(0,16000,100))
> D1log = lm(log(Time)~log(Tonnage)+I(log(Tonnage)^2),data=cargoes)
> CIs = exp(predict(D1log,new,interval="confidence"))
> PIs = exp(predict(D1log,new,interval="predict"))
> matpoints(new$Tonnage,CIs,lty=c(1,2,2),col=c("black","red","red"),type="l")
> matpoints(new$Tonnage,PIs,lty=c(1,2,2),col=c("black","blue","blue"),type="l")
```

## Simple Linear Regression

Transforming  $Y$  and/or  $X$  variables ("O, that way madness lies" King Lear)

- That's what we do

- ① Increasing variance: Transform  $Y$

The Hierarchy of transforms:

- square root – balanced but nasty
    - log – assertive but nice
    - inverse – aggressive and obnoxious

- ② Bad leverage points: Transform  $X$  (also see 5.)

- ③ Skewness in Standardized Residuals: Transform  $Y$

Often accompanies 1. together with presence of outliers

- ④ Skewness in  $X$  values: Transform  $X$

Often accompanies 2.

- ⑤ Incorrect functional form: Transform  $Y$  and/or  $X$  or add polynomial terms to the model